
Convolutional Neural Network (CNN)

卷积神经网络

Network Architecture designed for Image

Image Classification

one-hot vector
向量的长度 = 类别数量



100 x 100



$$\begin{bmatrix} \vdots \\ 0.2 \\ 0.7 \\ 0.1 \\ \vdots \end{bmatrix}$$

y'
prediction

dog $\begin{bmatrix} \vdots \\ 0 \end{bmatrix}$
cat $\begin{bmatrix} \vdots \\ 1 \end{bmatrix}$
tree $\begin{bmatrix} \vdots \\ 0 \end{bmatrix}$

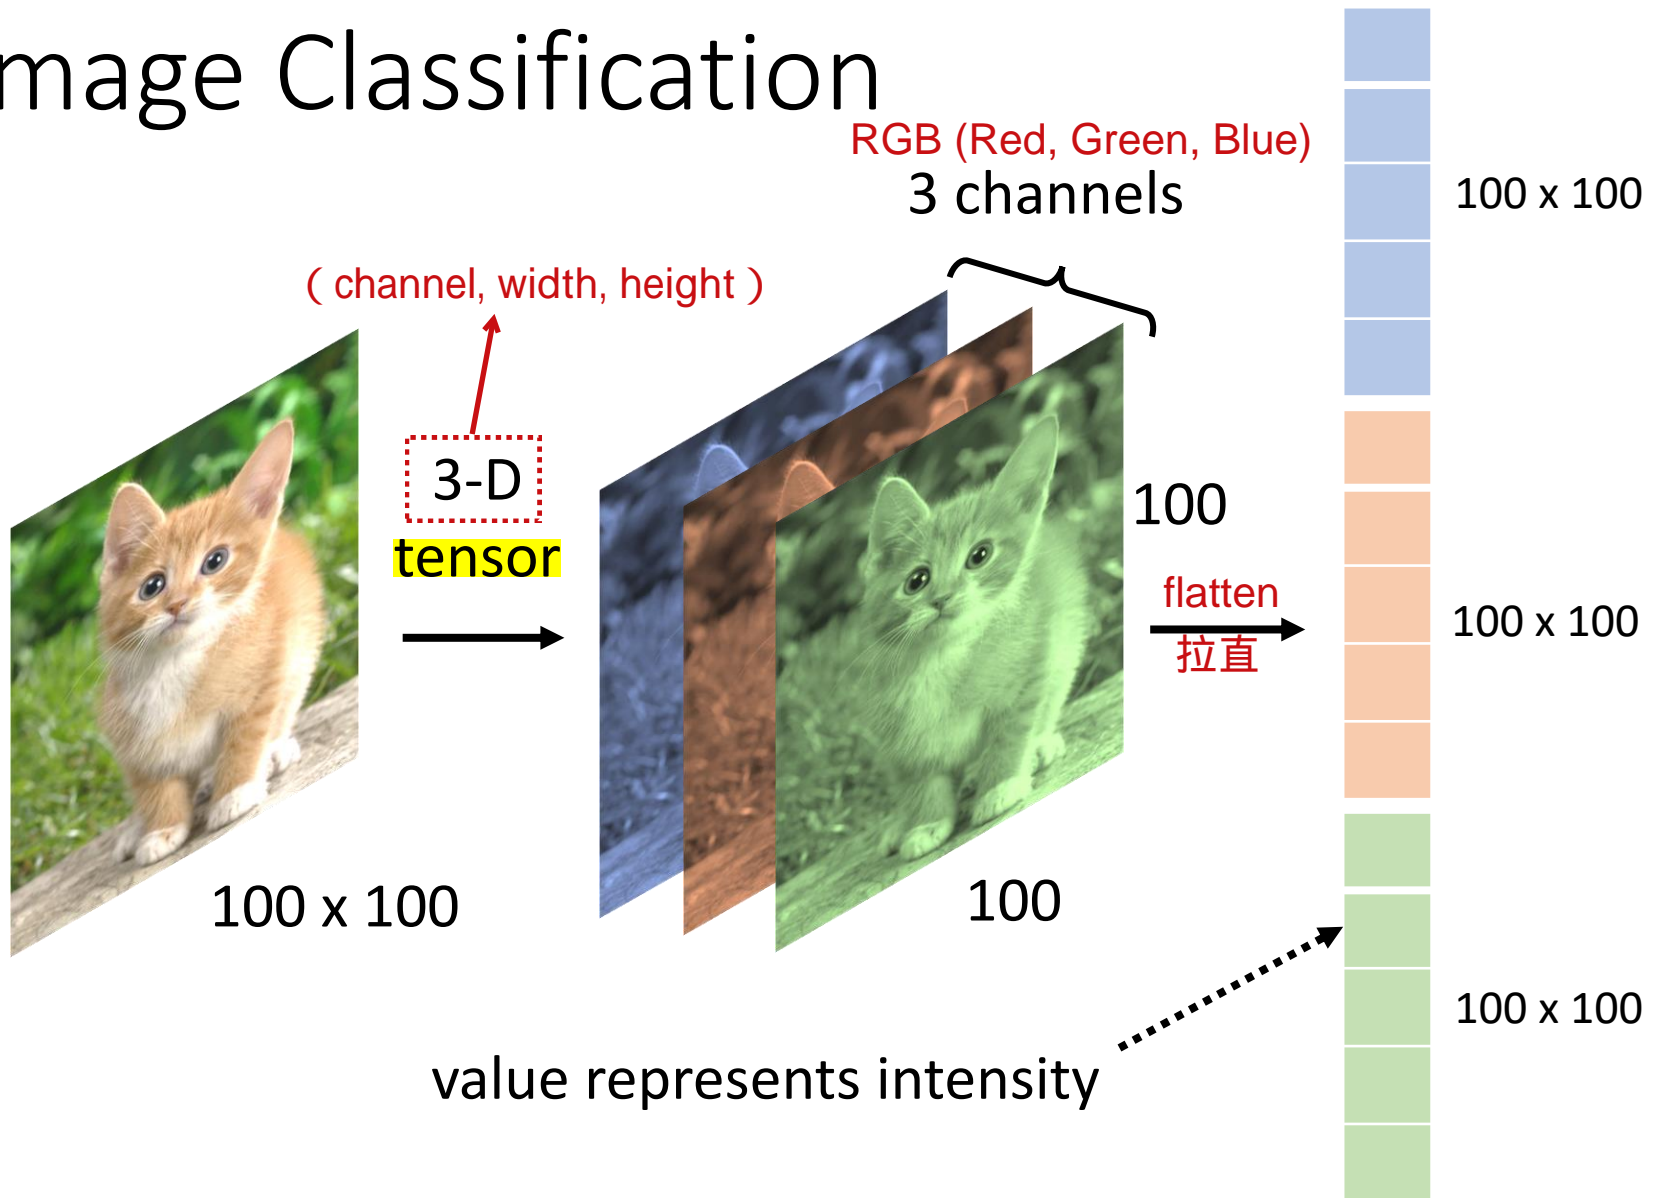
\hat{y}
target

Cross
entropy

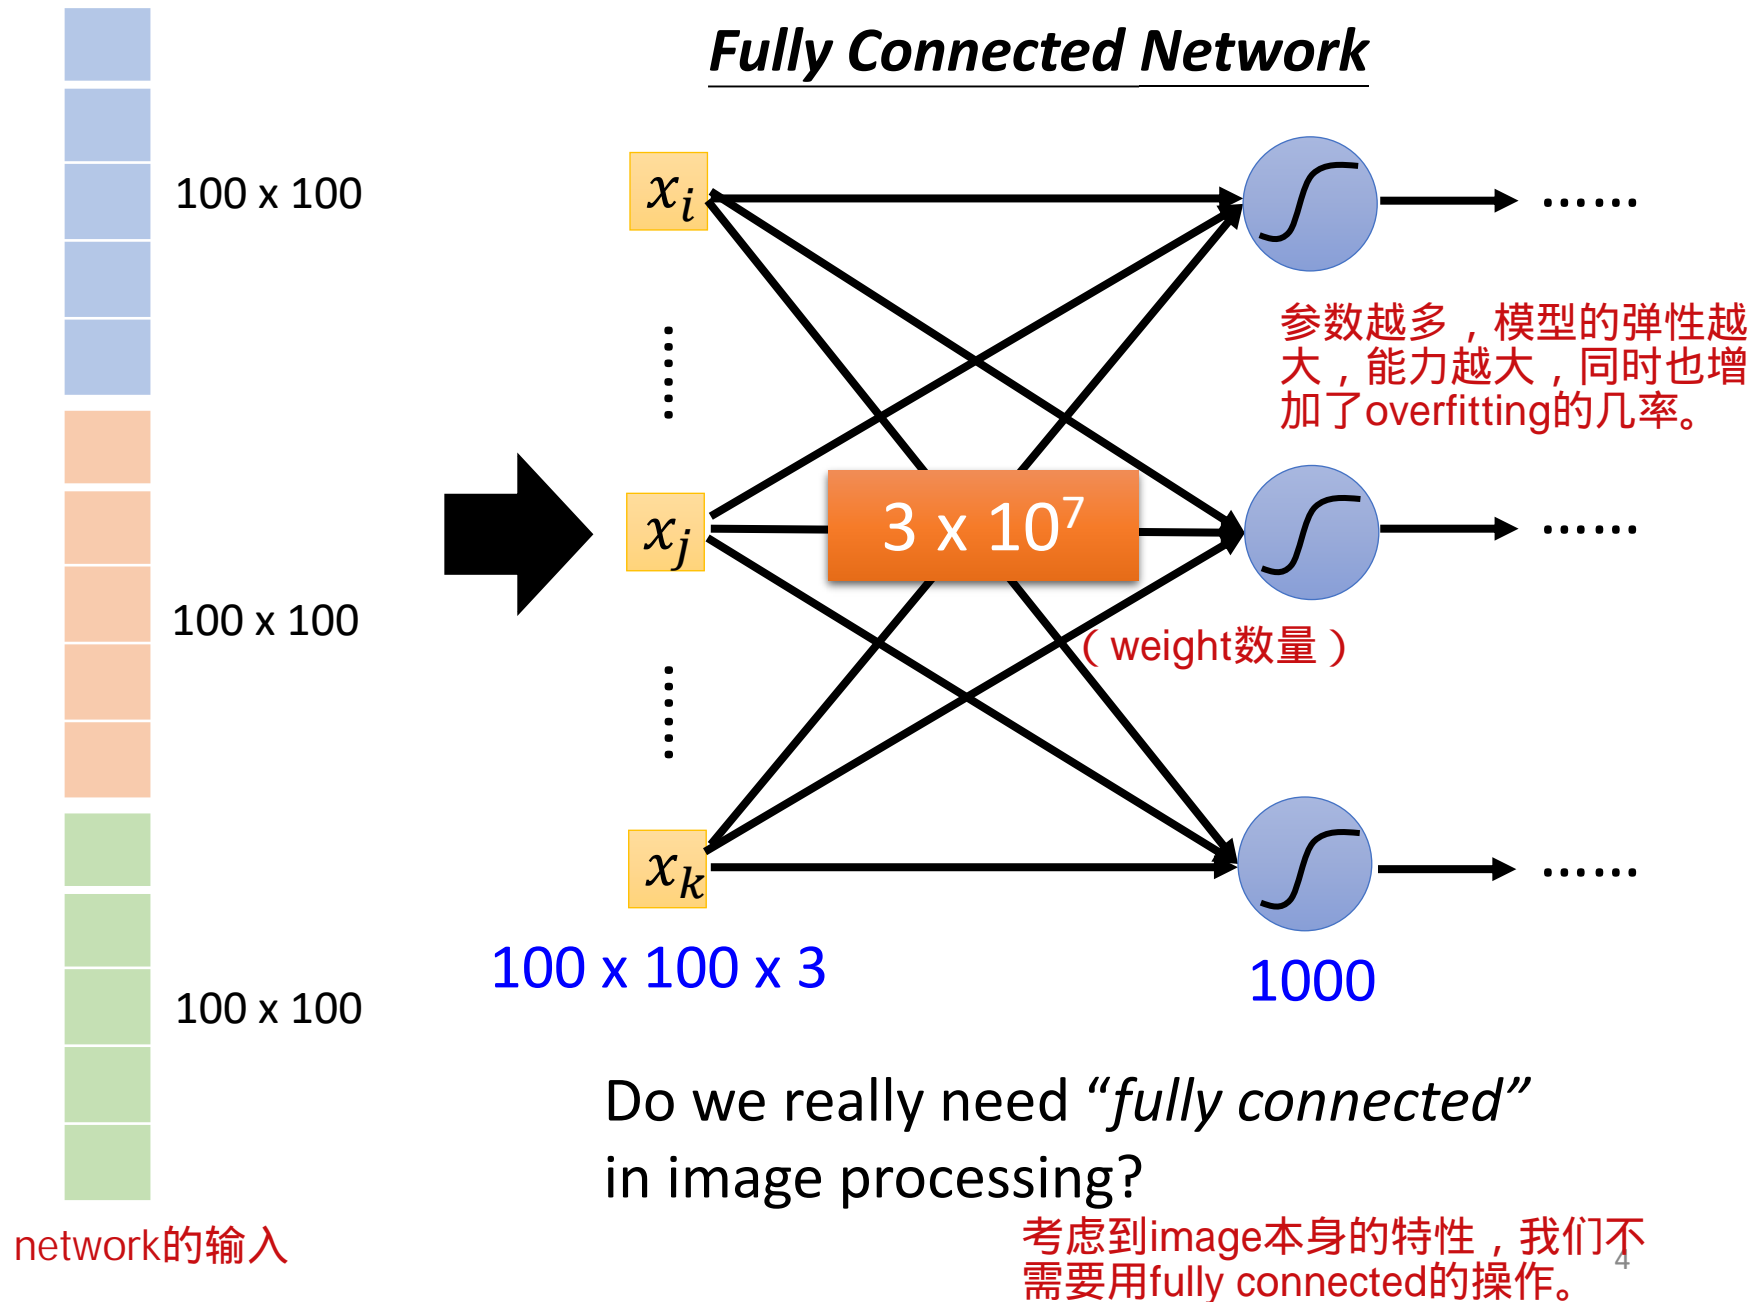
(All the images to be classified have the same size.)

输入图片的大小都要一样

Image Classification




Fully Connected Network



观察到某个重要的特性来做判断

Observation 1

Identifying some critical patterns



Perhaps human also identify birds in a similar way ... ☺

人也是抓住某个事物的特性来判断。



鸟的眼睛

鸟嘴

全身都是黑的


所有特征加起来，让有的人第一眼会觉得：这是乌鸦。

<https://www.dcard.tw/f/funny/p/233833012>

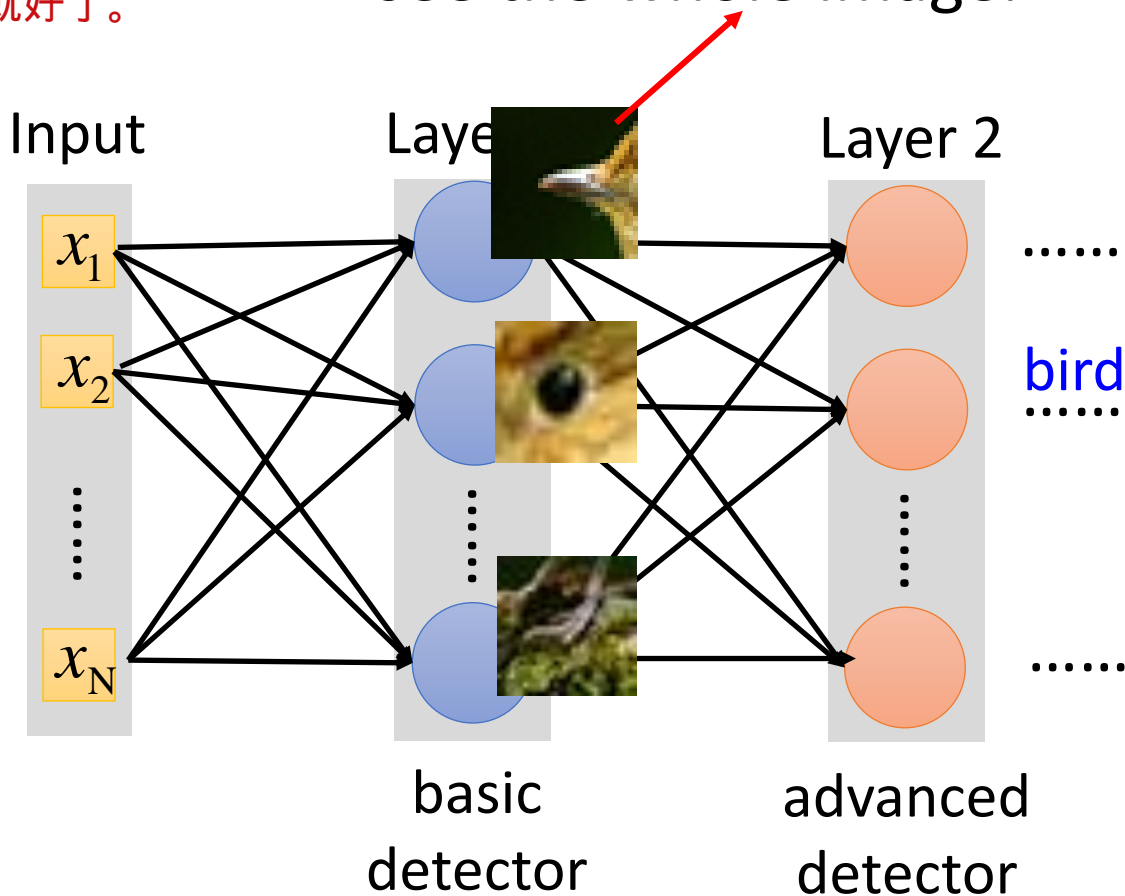
Observation 1

不需要看整张图片，看局部就好了。
不需要将整个图片作为输入。

Need to see the whole image?

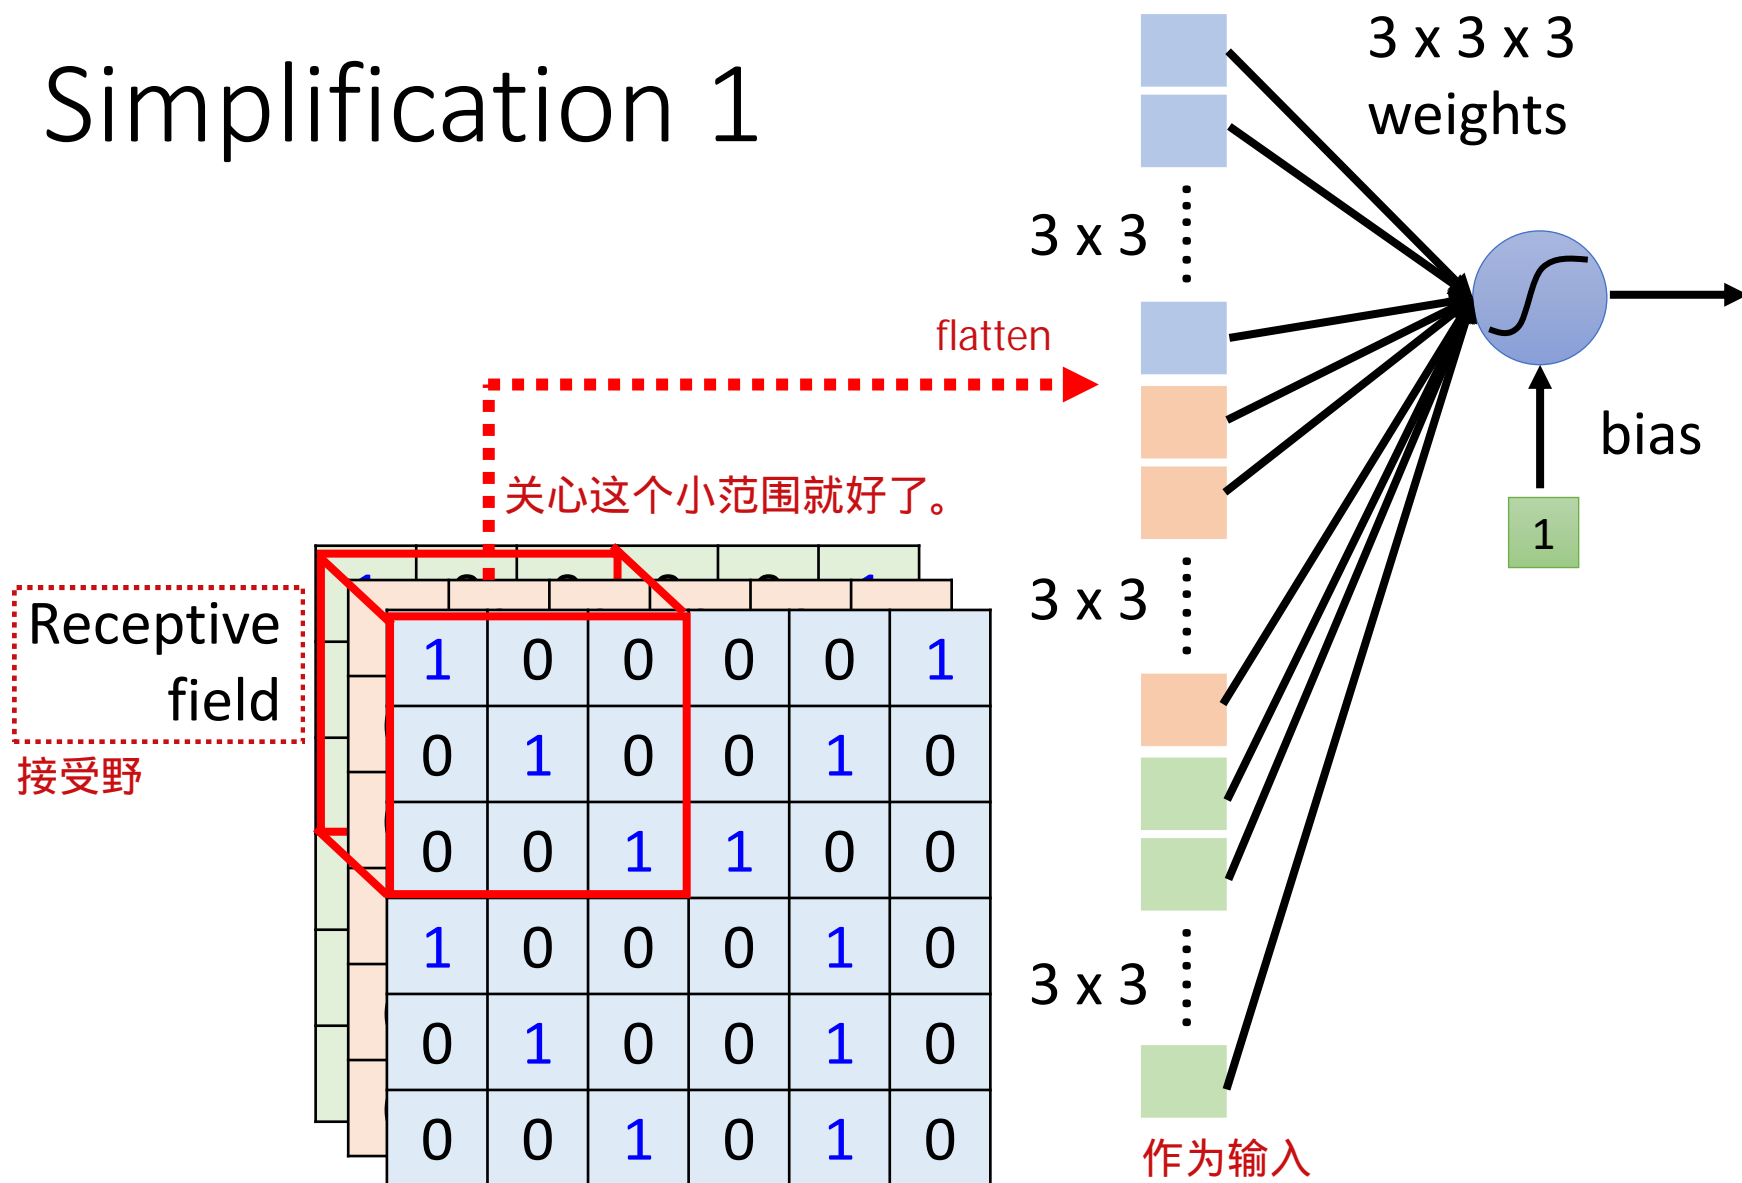


A neuron does not have to see the whole image.



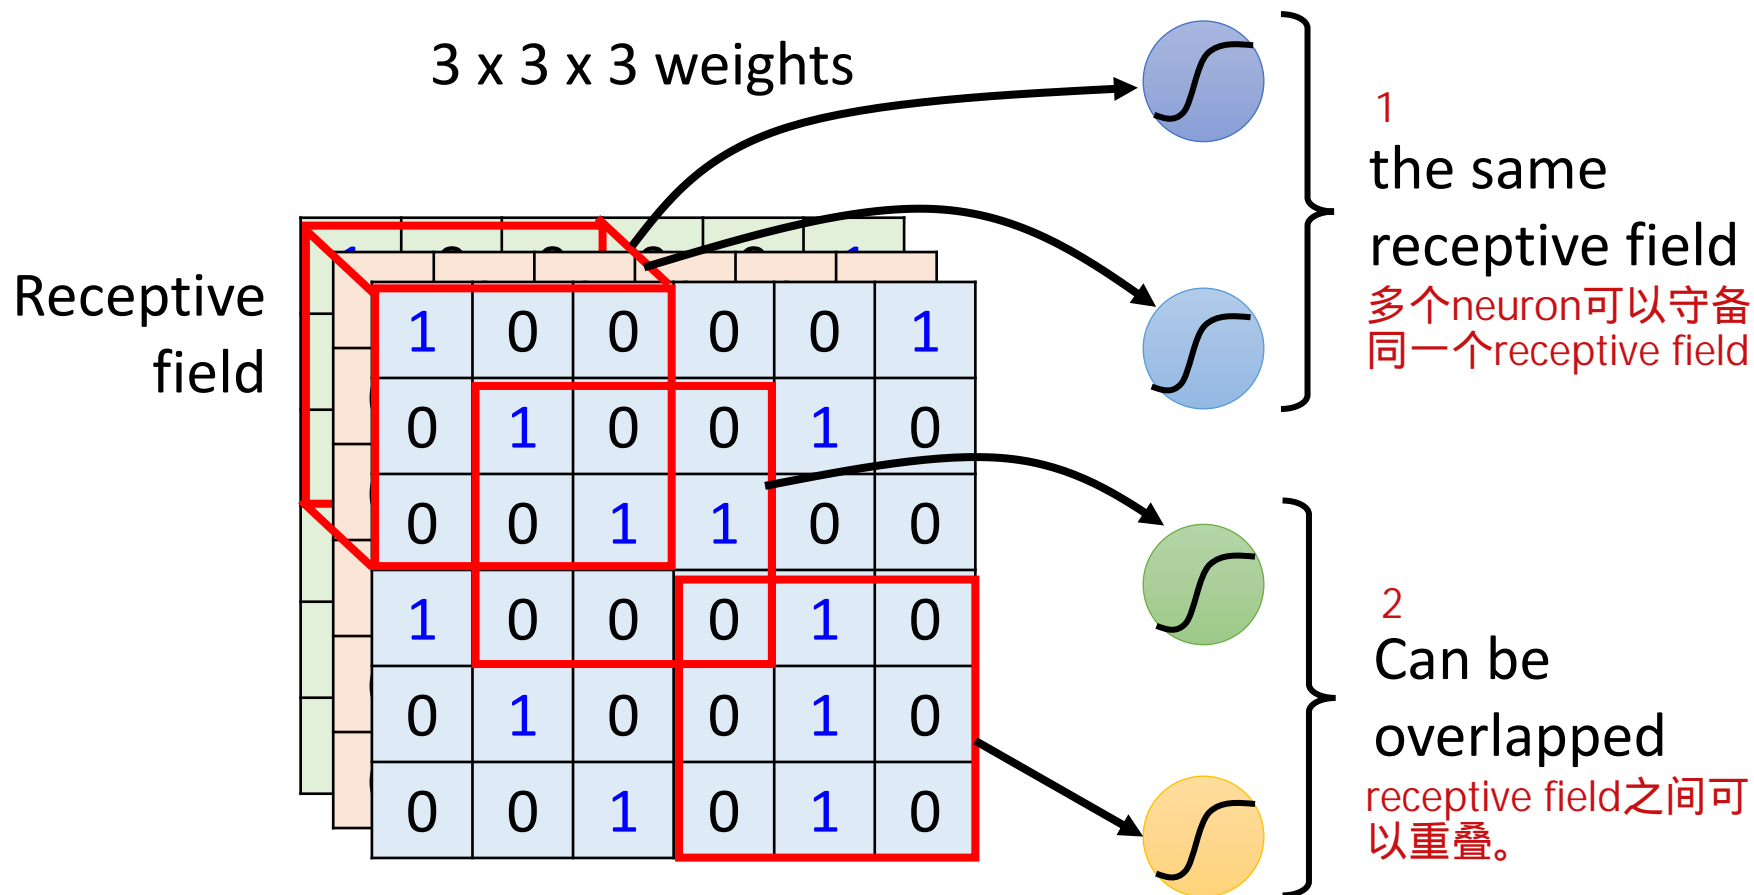
Some patterns are much smaller than the whole image.

Simplification 1



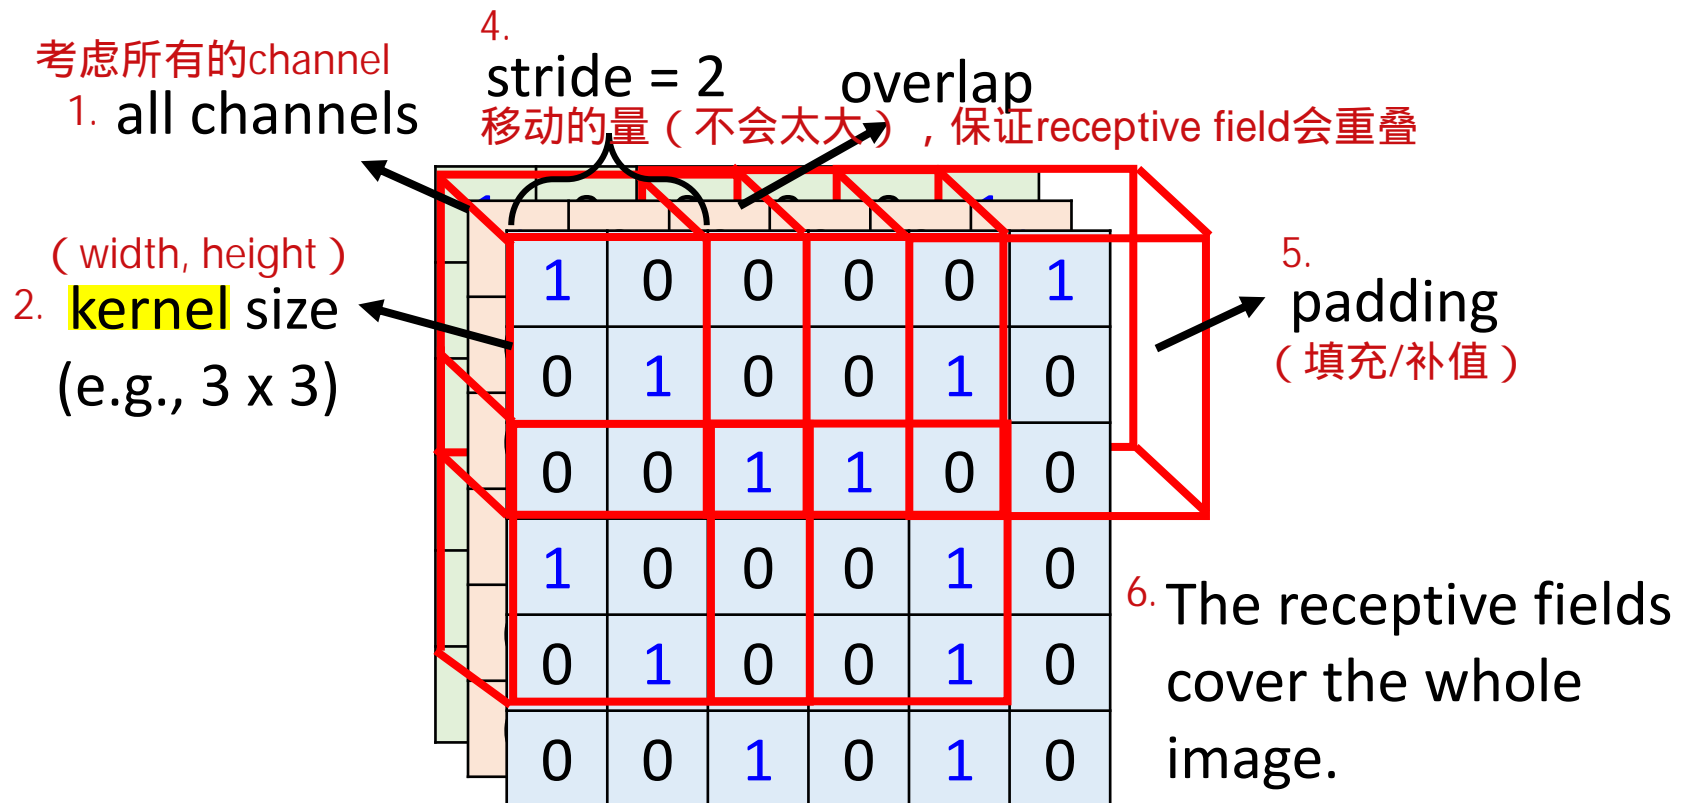
Simplification 1

- Can different neurons have different sizes of receptive field?
可以，常见招式，选择不同大小的receptive field。
- Cover only some channels?
可以，在network compression里会讲。
- Not square receptive field?
可以的，自己决定。



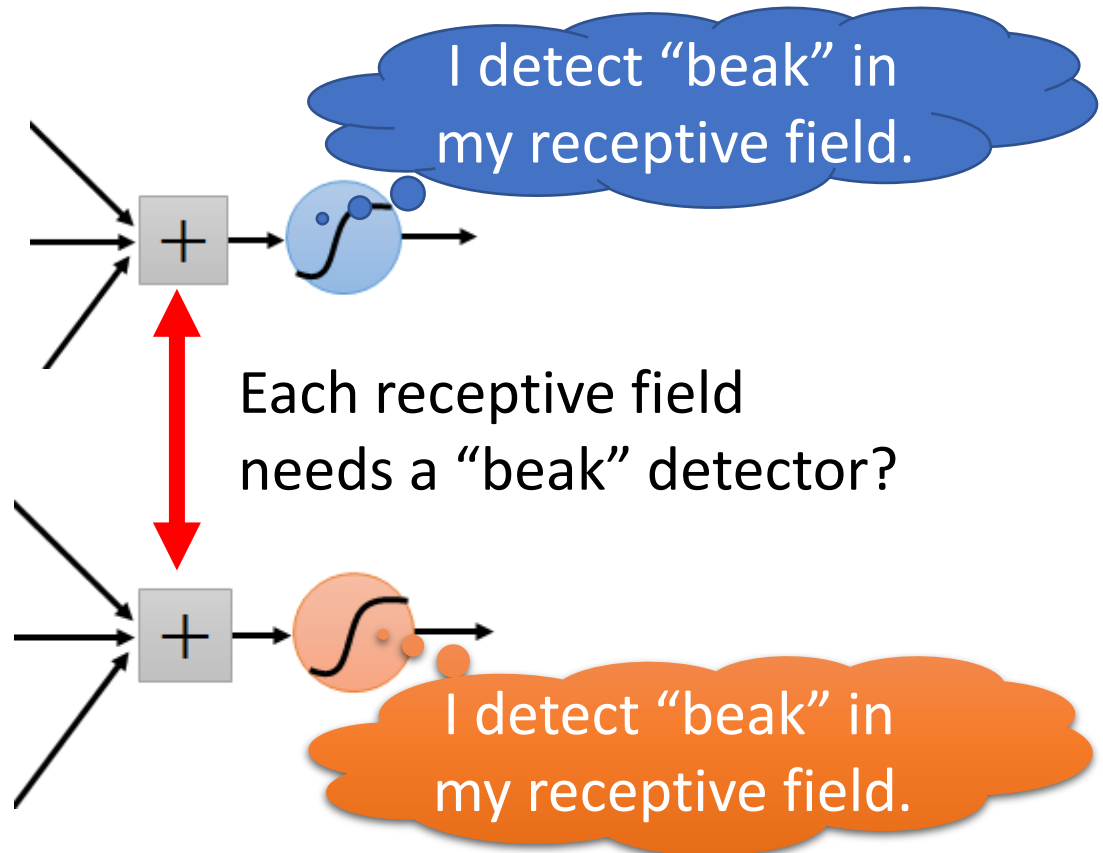
Simplification 1 – Typical Setting

3. Each receptive field has a set of neurons (e.g., 64 neurons).





Observation 2 特征在不同的区域

- The same patterns appear in different regions.




特征在不同的区域
solution : parameter sharing

Simplification 2





相同receptive field 不共享参数 不同output
不同receptive field 共享参数 不同output

Simplification 2




Two neurons with the same receptive field would not share parameters.



Simplification 2 – Typical Setting


Each receptive field has a set of neurons (e.g., 64 neurons).



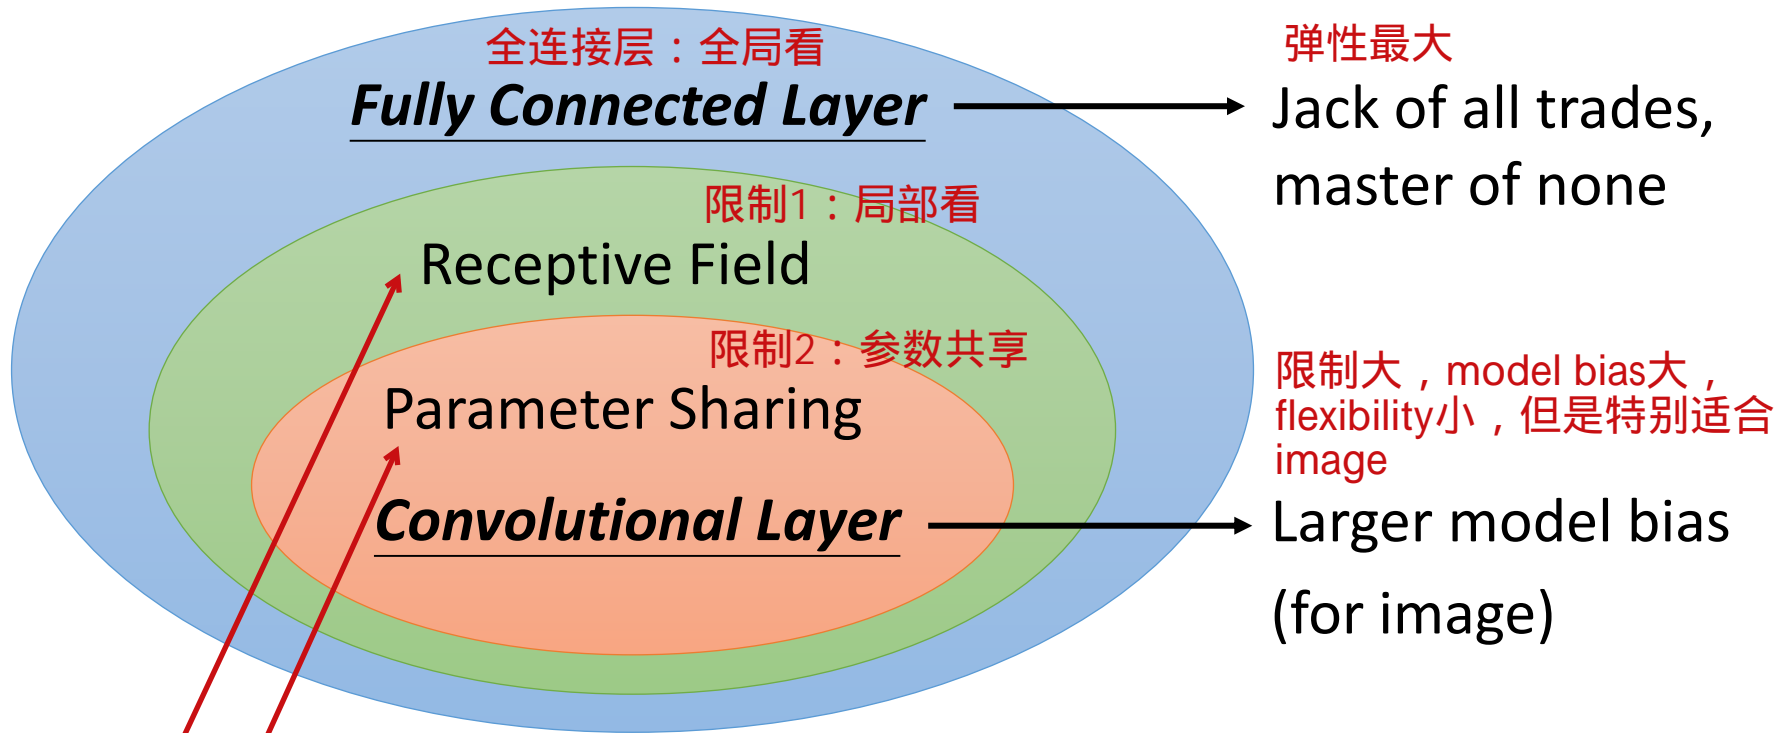
Simplification 2 – Typical Setting

Each receptive field has a set of neurons (e.g., 64 neurons).

Each receptive field has the neurons with the same set of parameters.



Benefit of Convolutional Layer




- Some patterns are much smaller than the whole image.
- The same patterns appear in different regions.

Another story based on **filter** 😊


Convolutional Layer



Convolution



Filter 1
3 x 3 x channel
tensor



Filter 2
3 x 3 x channel
tensor

channel = 3 (colorful) 彩色
channel = 1 (black and white) 黑白

Each filter detects a small pattern (3 x 3 x channel).

Convolutional Layer

Consider channel = 1
(black and white image)

1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0

6 x 6 image

1	-1	-1
-1	1	-1
-1	-1	1

Filter 1

-1	1	-1
-1	1	-1
-1	1	-1

Filter 2

⋮

(The values in the filters
are unknown parameters.)

Convolutional Layer

1	-1	-1
-1	1	-1
-1	-1	1

Filter 1

Inner Product

stride=1

1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0

6 x 6 image

3	-1	-3	-1
-3	1	0	-3
-3	-3	0	1
3	-2	-2	-1

左上角和左下角的值最大 有filter的特征

feature map

Convolutional Layer

-1	1	-1
-1	1	-1
-1	1	-1


Filter 2

stride=1

1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0

6 x 6 image

Do the same process for every filter



Convolutional Layer




64
filters

Convolution

Convolution

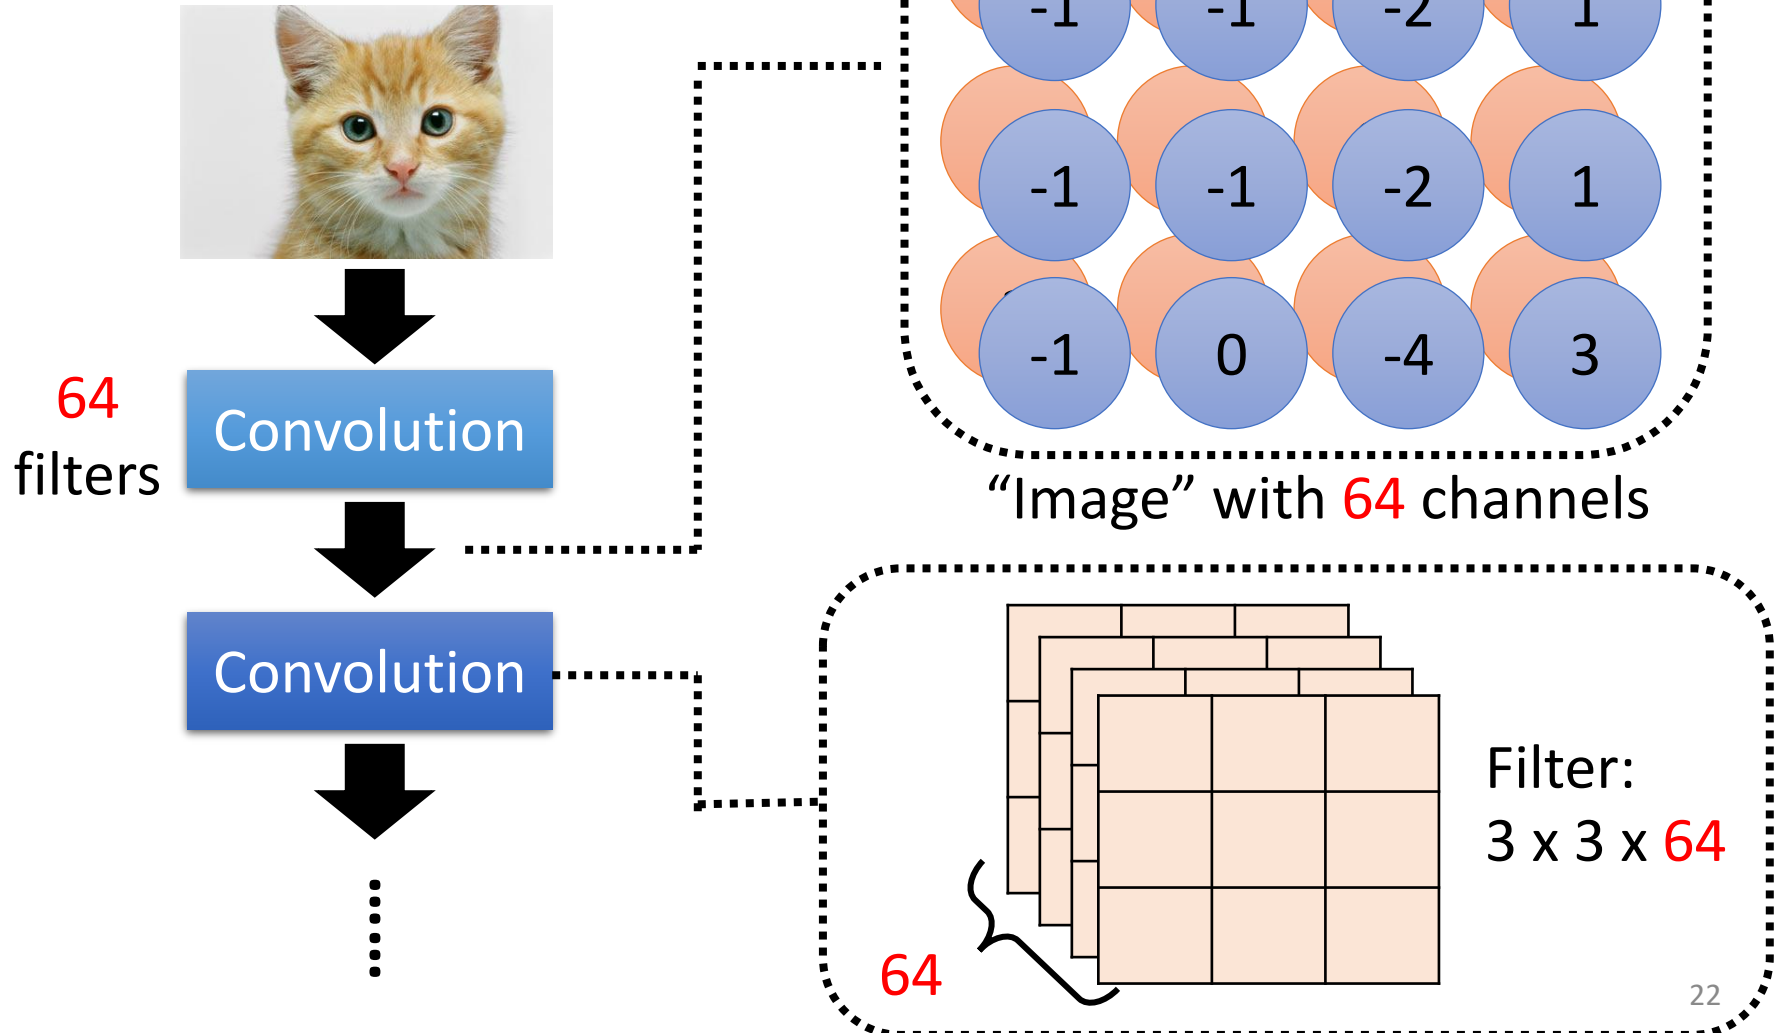
⋮



"Image" with 64 channels

可以看成是另外一张图片，但是是64个channels。
of channels = # of filters of the last layer

Multiple Convolutional Layers



Multiple Convolutional Layers



64
filters

Convolution


Convolution

⋮

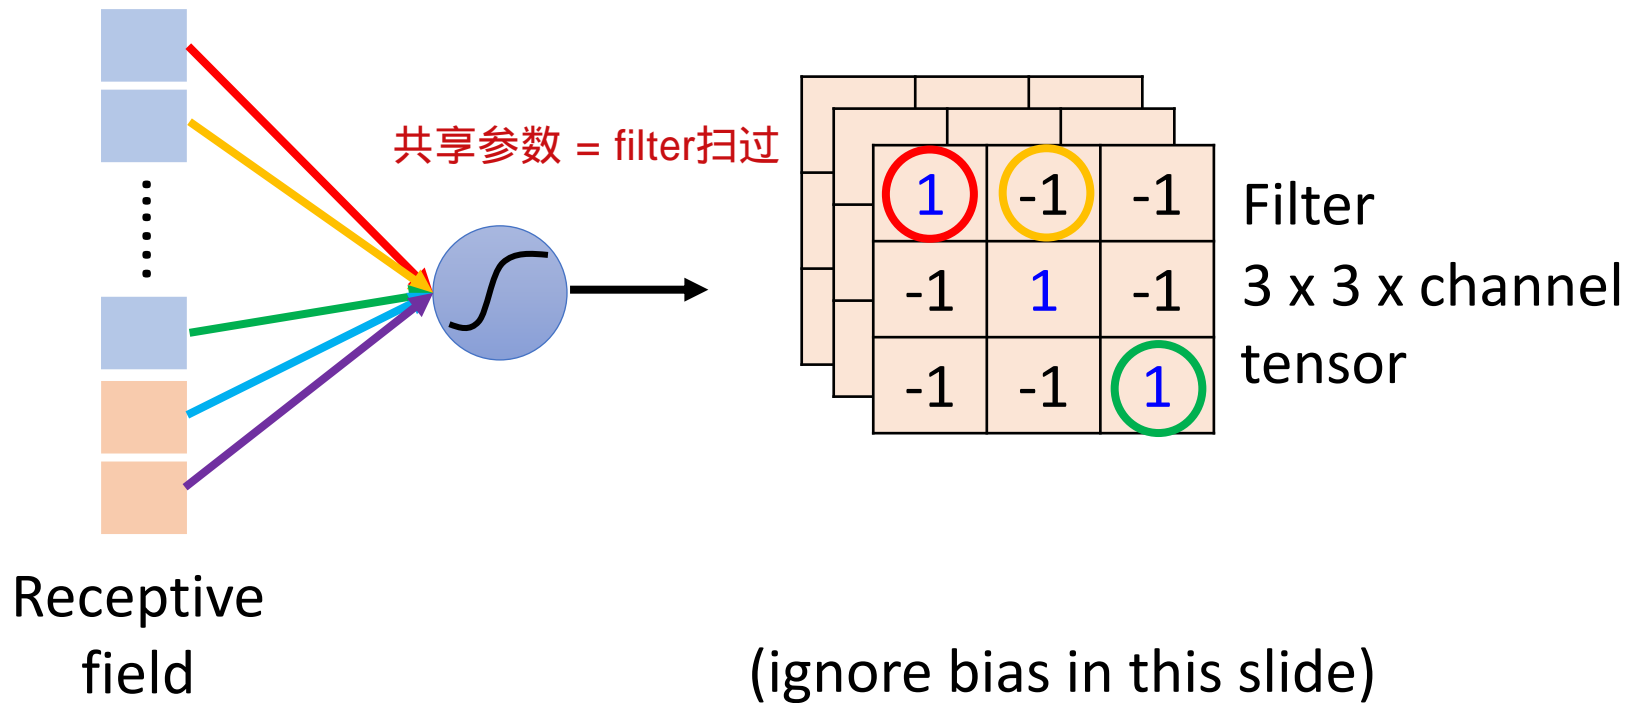
实际上考虑的是5×5

1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0

看起来是3×3




Comparison of Two Stories




convolution = filter扫过整张图片 = receptive field共享参数

The neurons with different receptive fields **share the parameters**.



Each filter convolves over the input image.



Convolutional Layer

<u>Neuron Version Story</u>	<u>Filter Version Story</u>
Each neuron only considers a receptive field.	There are a set of filters detecting small patterns.
The neurons with different receptive fields share the parameters.	Each filter convolves over the input image.

They are the same story.

pooling

Observation 3

- Subsampling the pixels will not change the object

bird



subsampling

下采样

bird



pooling不会学习；没有参数

Pooling – Max Pooling

选的代表：最大的那个



有很多pooling的方法；
e.g., mean pooling

1	-1	-1
-1	1	-1
-1	-1	1


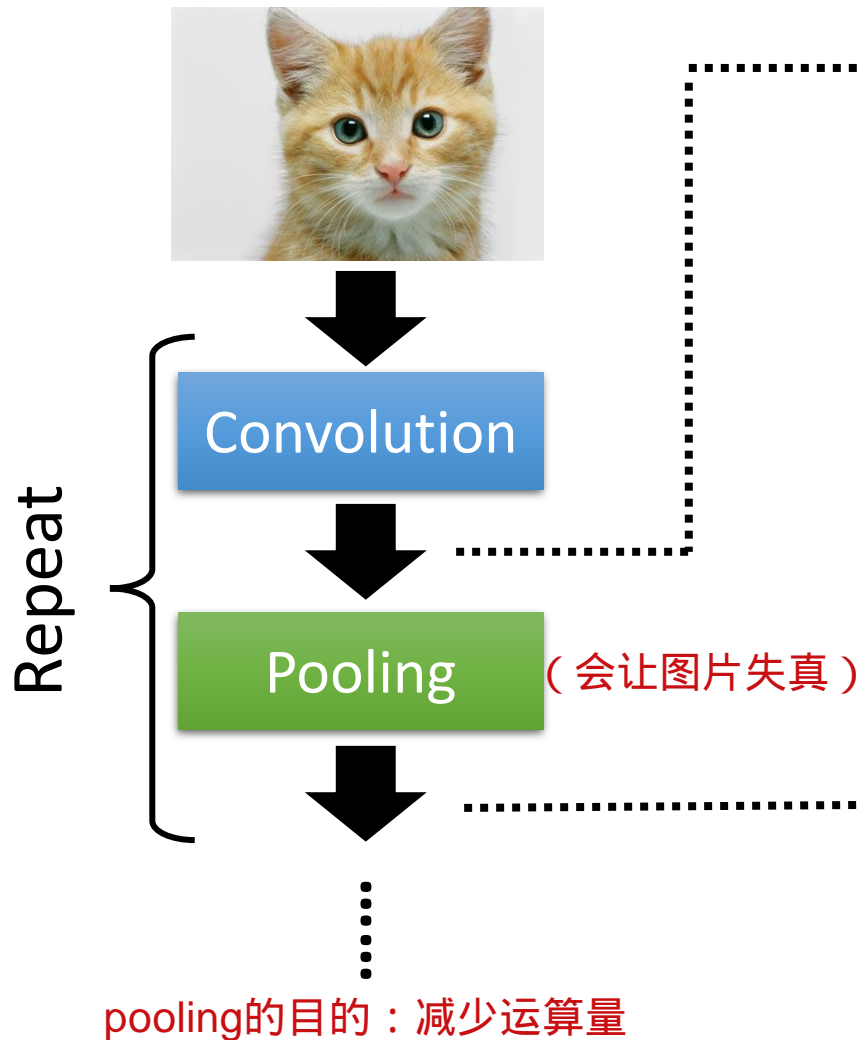
Filter 1

-1	1	-1
-1	1	-1
-1	1	-1


Filter 2



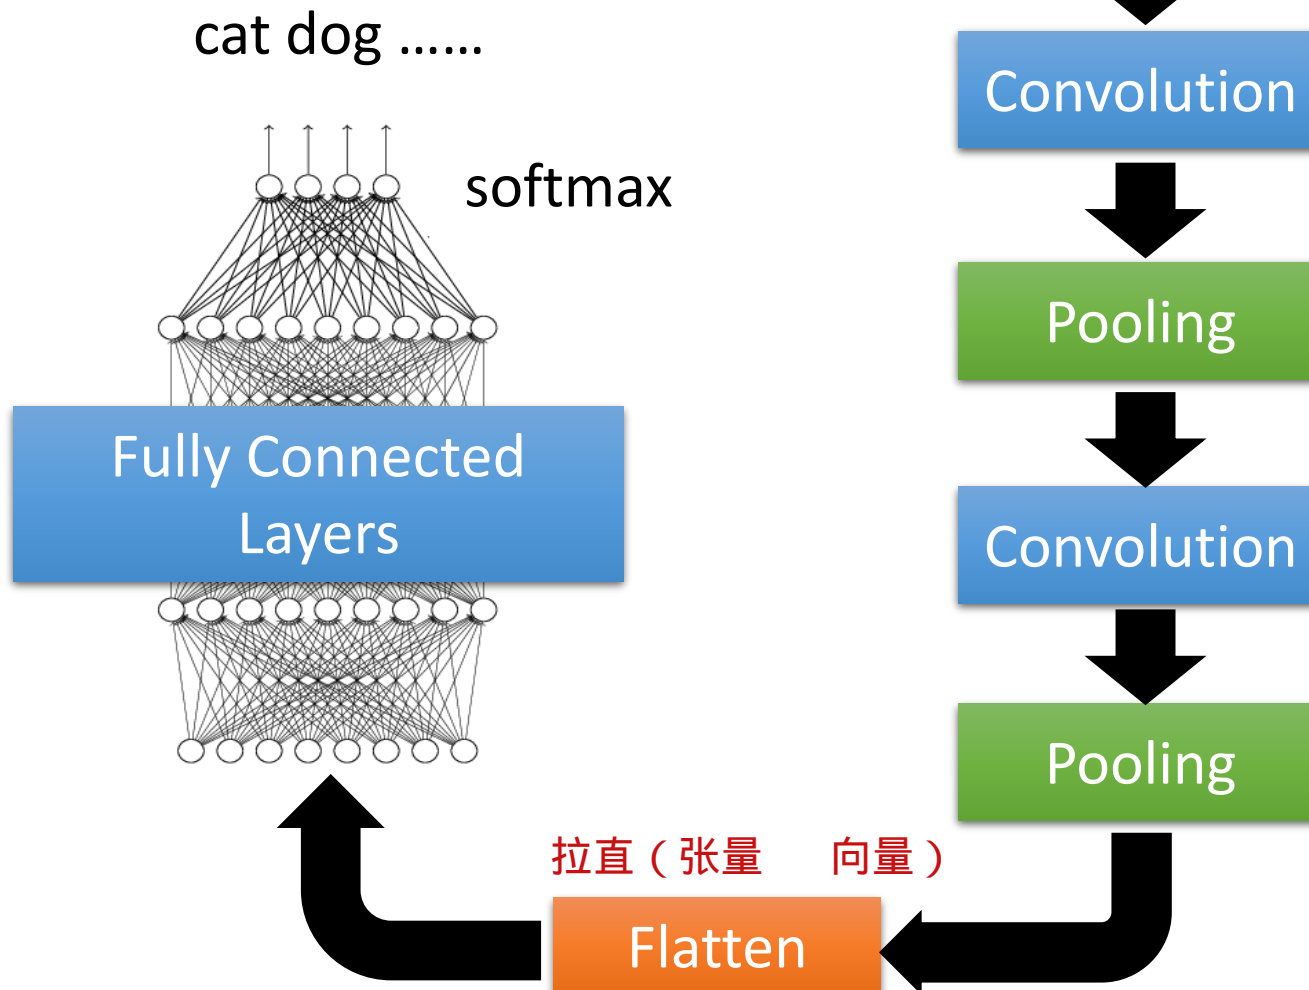
Convolutional Layers + Pooling



"Image" with 64 channels



The whole CNN



Application: Playing Go 下围棋 = 分类问题



Next move
(19 x 19 小的图片
positions)

19 x 19 classes
(19, 19, 48)

domain knowledge
48 channels
in Alpha Go

Black: 1
white: -1
none: 0

Fully-connected
network can be used


But CNN performs much better.

当任务和图像处理有类似特性时，可以用CNN。



Why CNN for Go playing?

- 1 • Some patterns are much smaller than the whole image

Alpha Go uses 5 x 5 for first layer



- 2 • The same patterns appear in different regions.



AlphaGo的CNN不用Pooling，
因为Pooling会掉信息，所以不
适合下棋的任务。

Why CNN for Go playing?

- Subsampling the pixels will not change the object



Pooling

How to explain this???

Neural network architecture. The input to the policy network is a $19 \times 19 \times 48$ image stack consisting of 48 feature planes. The first hidden layer zero pads the input into a 23×23 image, then convolves k filters of kernel size 5×5 with stride 1 with the input image and applies a ^{ReLU} rectifier nonlinearity. Each of the subsequent hidden layers 2 to 12 zero pads the respective previous hidden layer into a 21×21 image, then convolves k filters of kernel size 3×3 with stride 1, again followed by a rectifier nonlinearity. The final layer convolves 1 filter of kernel size 1×1 with stride 1, with a different bias for each position, and applies a softmax function. The match version of AlphaGo used $k = 192$ filters; Fig. 2b and Extended Data Table 1. The tournament version used 256 and 384 filters.

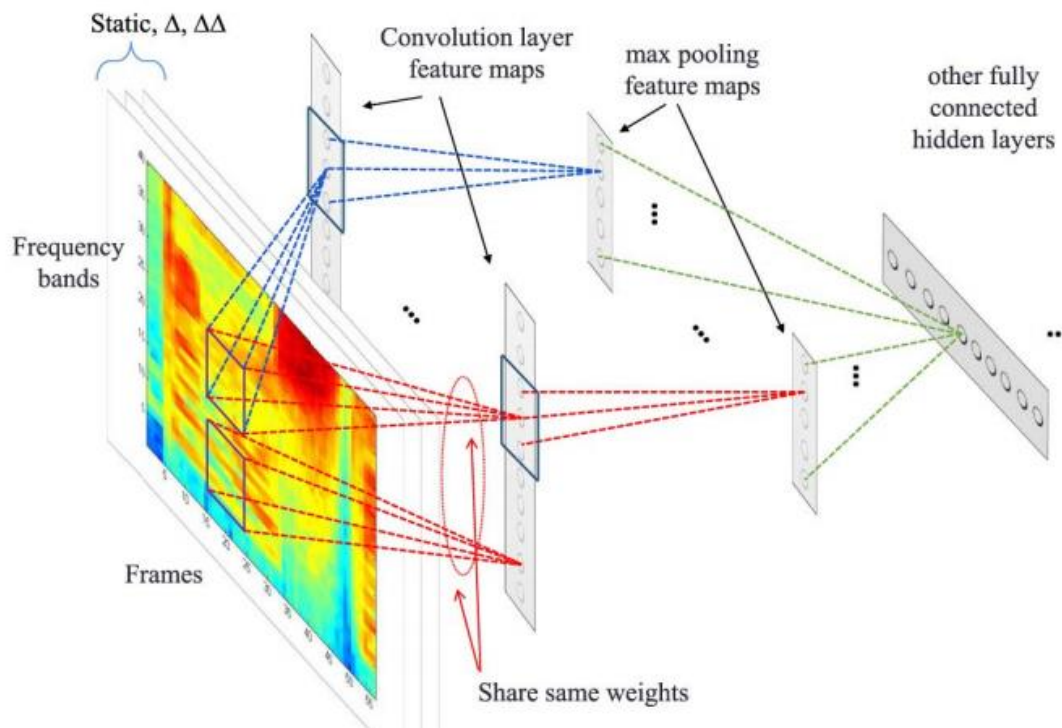
Alpha Go does not use Pooling

More Applications

语音和NLP上也都有应用。

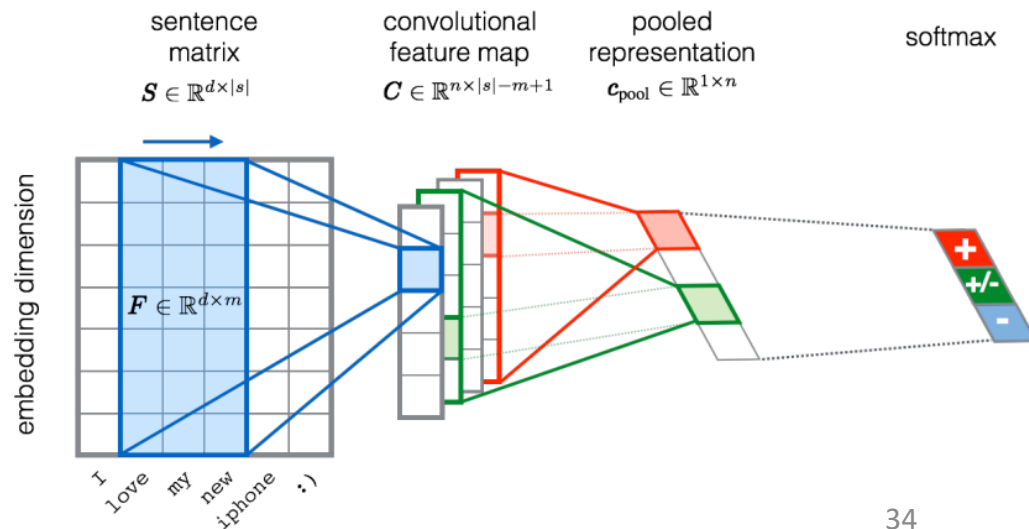
Speech

<https://dl.acm.org/doi/10.1109/TASLP.2014.2339736>



Natural Language Processing

<https://www.aclweb.org/anthology/S15-2079/>



To learn more ...

- CNN is not invariant to 放大/缩小 scaling and 旋转 rotation (we need data augmentation 😊).



Spatial Transformer Layer



<https://youtu.be/SoCywZ1hZak>
(in Mandarin)