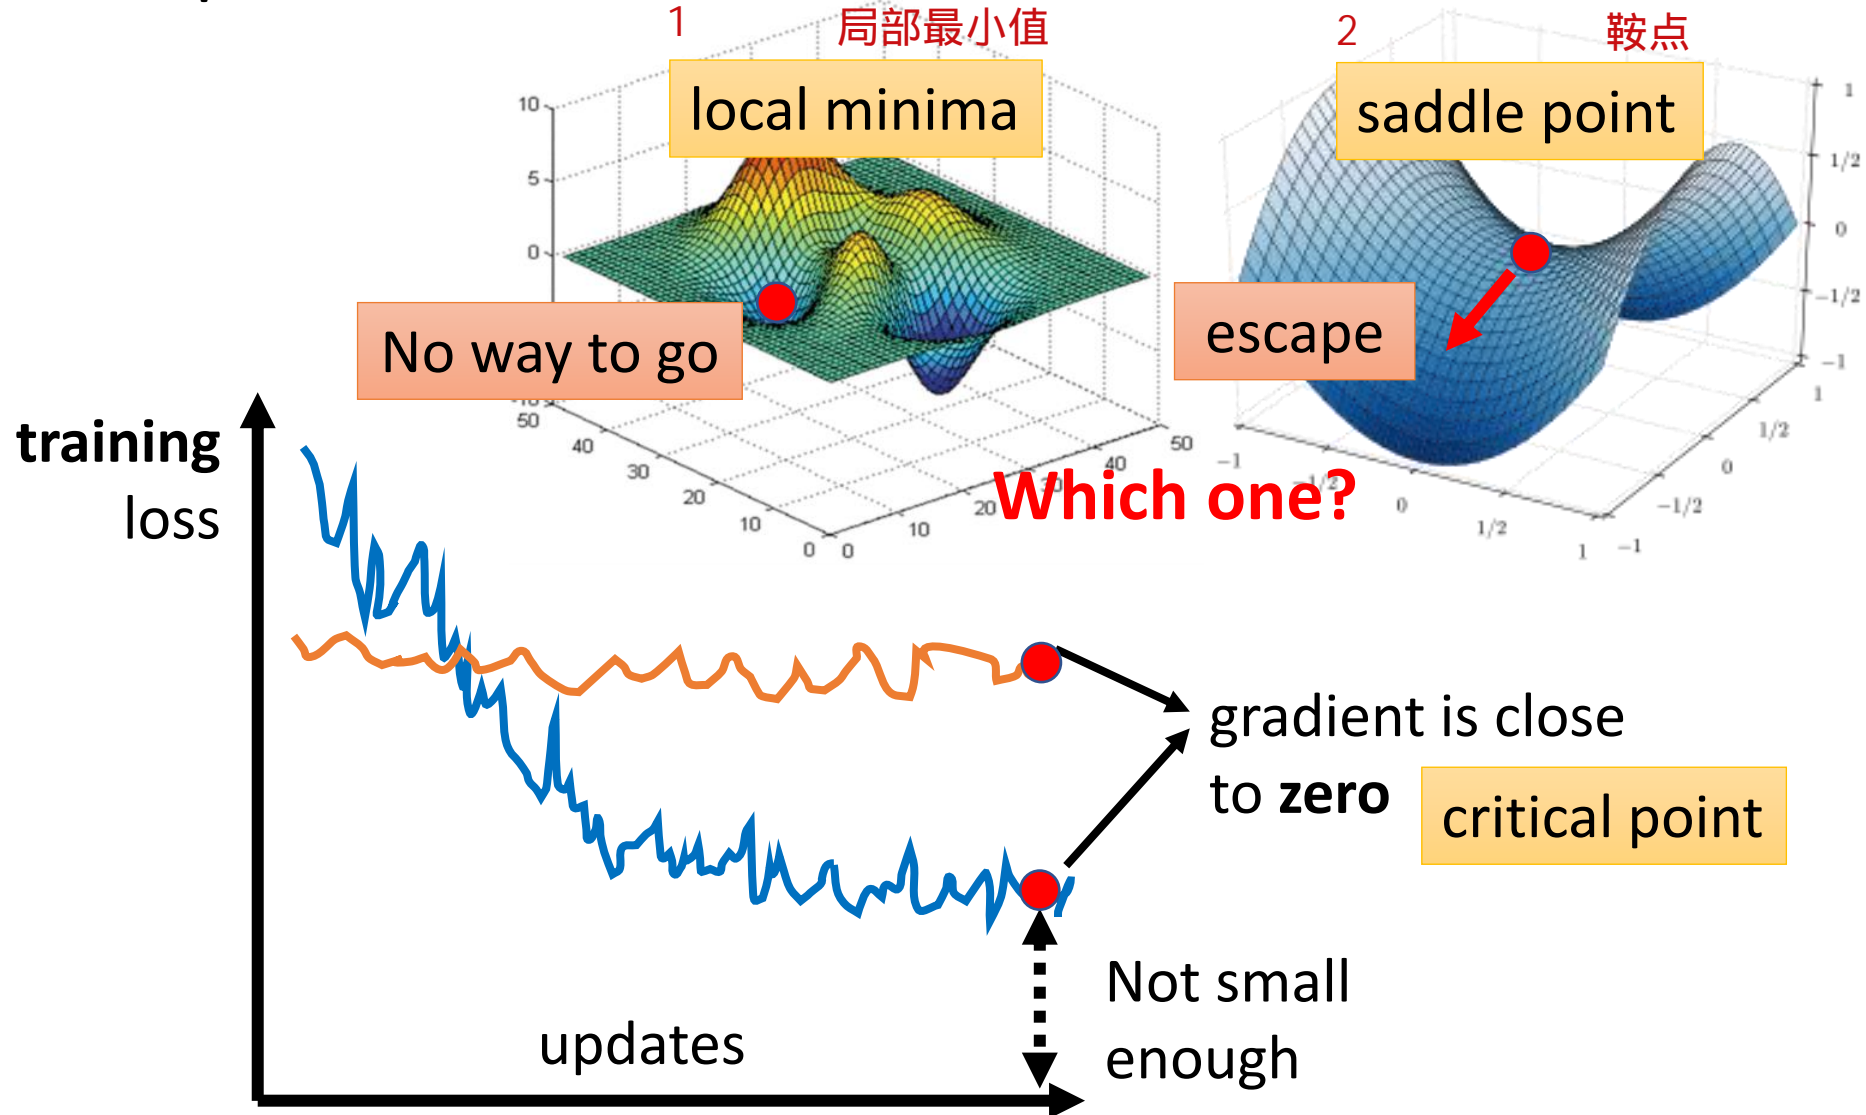




When gradient is small ...

Hung-yi Lee 李宏毅

Optimization Fails because



Warning of Math

你过来啊！

Taylor Series Approximation

泰勒级数展开

$L(\boldsymbol{\theta})$ around $\boldsymbol{\theta} = \boldsymbol{\theta}'$ can be approximated below


$$L(\boldsymbol{\theta}) \approx L(\boldsymbol{\theta}') + (\boldsymbol{\theta} - \boldsymbol{\theta}')^T \boldsymbol{g} + \frac{1}{2} (\boldsymbol{\theta} - \boldsymbol{\theta}')^T \boldsymbol{H} (\boldsymbol{\theta} - \boldsymbol{\theta}')$$

Gradient \boldsymbol{g} is a vector

$$\boldsymbol{g} = \nabla L(\boldsymbol{\theta}') \quad g_i = \frac{\partial L(\boldsymbol{\theta}')}{\partial \theta_i}$$

Hessian \boldsymbol{H} is a matrix

$$H_{ij} = \frac{\partial^2}{\partial \theta_i \partial \theta_j} L(\boldsymbol{\theta}')$$



Hessian


$L(\boldsymbol{\theta})$ around $\boldsymbol{\theta} = \boldsymbol{\theta}'$ can be approximated below

$$L(\boldsymbol{\theta}) \approx L(\boldsymbol{\theta}') + \cancel{(\boldsymbol{\theta} - \boldsymbol{\theta}')^T \mathbf{g}} + \frac{1}{2} (\boldsymbol{\theta} - \boldsymbol{\theta}')^T \mathbf{H} (\boldsymbol{\theta} - \boldsymbol{\theta}')$$


At critical point

telling the properties of critical points


local min



local max



saddle point



算出Hessian矩阵，计算所有特征值：

1. all positive local minima
2. all negative local maxima
3. else saddle point

Hessian

At critical point:

$$\mathbf{v}^T \mathbf{H} \mathbf{v}$$

$$L(\boldsymbol{\theta}) \approx L(\boldsymbol{\theta}') + \frac{1}{2} (\boldsymbol{\theta} - \boldsymbol{\theta}')^T \mathbf{H} (\boldsymbol{\theta} - \boldsymbol{\theta}')$$

For all \mathbf{v}


1. $\mathbf{v}^T \mathbf{H} \mathbf{v} > 0 \implies$ Around $\boldsymbol{\theta}'$: $L(\boldsymbol{\theta}) > L(\boldsymbol{\theta}')$ \implies Local minima
= \mathbf{H} is positive definite = All eigen values are positive. \uparrow
特征值

For all \mathbf{v}


2. $\mathbf{v}^T \mathbf{H} \mathbf{v} < 0 \implies$ Around $\boldsymbol{\theta}'$: $L(\boldsymbol{\theta}) < L(\boldsymbol{\theta}')$ \implies Local maxima
= \mathbf{H} is negative definite = All eigen values are negative. \uparrow
3. Sometimes $\mathbf{v}^T \mathbf{H} \mathbf{v} > 0$, sometimes $\mathbf{v}^T \mathbf{H} \mathbf{v} < 0 \implies$ Saddle point
Some eigen values are positive, and some are negative. \uparrow



Example

$$y = w_1 w_2 x \quad (\text{史上最废neural network})$$



Error Surface





$$L = (\hat{y} - w_1 w_2 x)^2 = (1 - w_1 w_2)^2$$

$$\frac{\partial L}{\partial w_1} = 2(1 - w_1 w_2)(-w_2) = 0$$

$$\frac{\partial L}{\partial w_2} = 2(1 - w_1 w_2)(-w_1) = 0$$

Critical point: $w_1 = 0, w_2 = 0$

$$H = \begin{bmatrix} 0 & -2 \\ -2 & 0 \end{bmatrix} \quad \lambda_1 = 2, \lambda_2 = -2$$

Saddle point

g

H

$$\frac{\partial^2 L}{\partial w_1^2} = 2(-w_2)(-w_2) = 0$$

$$\frac{\partial^2 L}{\partial w_2 \partial w_1} = -2 + 4w_1 w_2 = -2$$

$$\frac{\partial^2 L}{\partial w_1 \partial w_2} = -2 + 4w_1 w_2 = -2$$

$$\frac{\partial^2 L}{\partial w_2^2} = 2(-w_1)(-w_1) = 0$$

Don't afraid of saddle point?

$$\mathbf{v}^T \mathbf{H} \mathbf{v}$$

At critical point: $L(\boldsymbol{\theta}) \approx L(\boldsymbol{\theta}') + \frac{1}{2} (\boldsymbol{\theta} - \boldsymbol{\theta}')^T \mathbf{H} (\boldsymbol{\theta} - \boldsymbol{\theta}')$

Sometimes $\mathbf{v}^T \mathbf{H} \mathbf{v} > 0$, sometimes $\mathbf{v}^T \mathbf{H} \mathbf{v} < 0 \Rightarrow$ Saddle point

\mathbf{H} may tell us parameter update direction!

\mathbf{u} is an eigen vector of \mathbf{H}

λ is the eigen value of \mathbf{u}

$$\lambda < 0$$





$$\mathbf{u}^T \mathbf{H} \mathbf{u} = \mathbf{u}^T (\lambda \mathbf{u}) = \lambda \|\mathbf{u}\|^2$$
$$< 0 \qquad \qquad \qquad < 0$$

$$L(\boldsymbol{\theta}) \approx L(\boldsymbol{\theta}') + \frac{1}{2} (\boldsymbol{\theta} - \boldsymbol{\theta}')^T \mathbf{H} (\boldsymbol{\theta} - \boldsymbol{\theta}') \Rightarrow L(\boldsymbol{\theta}) < L(\boldsymbol{\theta}')$$

$$\boldsymbol{\theta} - \boldsymbol{\theta}' = \mathbf{u}$$

$$\boldsymbol{\theta} = \boldsymbol{\theta}' + \mathbf{u}$$

Decrease L



$$L = (\hat{y} - w_1 w_2 x)^2 = (1 - w_1 w_2)^2$$

$$\frac{\partial L}{\partial w_1} = 2(1 - w_1 w_2)(-w_2)$$

$$\frac{\partial L}{\partial w_2} = 2(1 - w_1 w_2)(-w_1)$$

Critical point: $w_1 = 0, w_2 = 0$

$$H = \begin{bmatrix} 0 & -2 \\ -2 & 0 \end{bmatrix} \quad \lambda_1 = 2, \lambda_2 = -2$$

Saddle point

$\lambda_2 = -2$ Has eigenvector $\mathbf{u} = \begin{bmatrix} 1 \\ 1 \end{bmatrix}$

Update the parameter along the direction of \mathbf{u}

You can escape the saddle point and decrease the loss.

实际中，不会用Hessian来逃离
saddle point，计算量太大啦

(this method is seldom used in practice)

End of Warning

Saddle Point v.s. Local Minima

- A.D. 1543

1453年
君士坦丁堡沦陷
Fall of Constantinople



Saddle Point v.s. Local Minima

- The Magician Diorena (魔法師狄奧倫娜)

From 3 dimensional space, it is sealed.


It is not in higher dimensions.





Source of image: <https://read01.com/mz2DBPE.html#.YECz22gzblU>

Saddle Point v.s. Local Minima

二维



三维



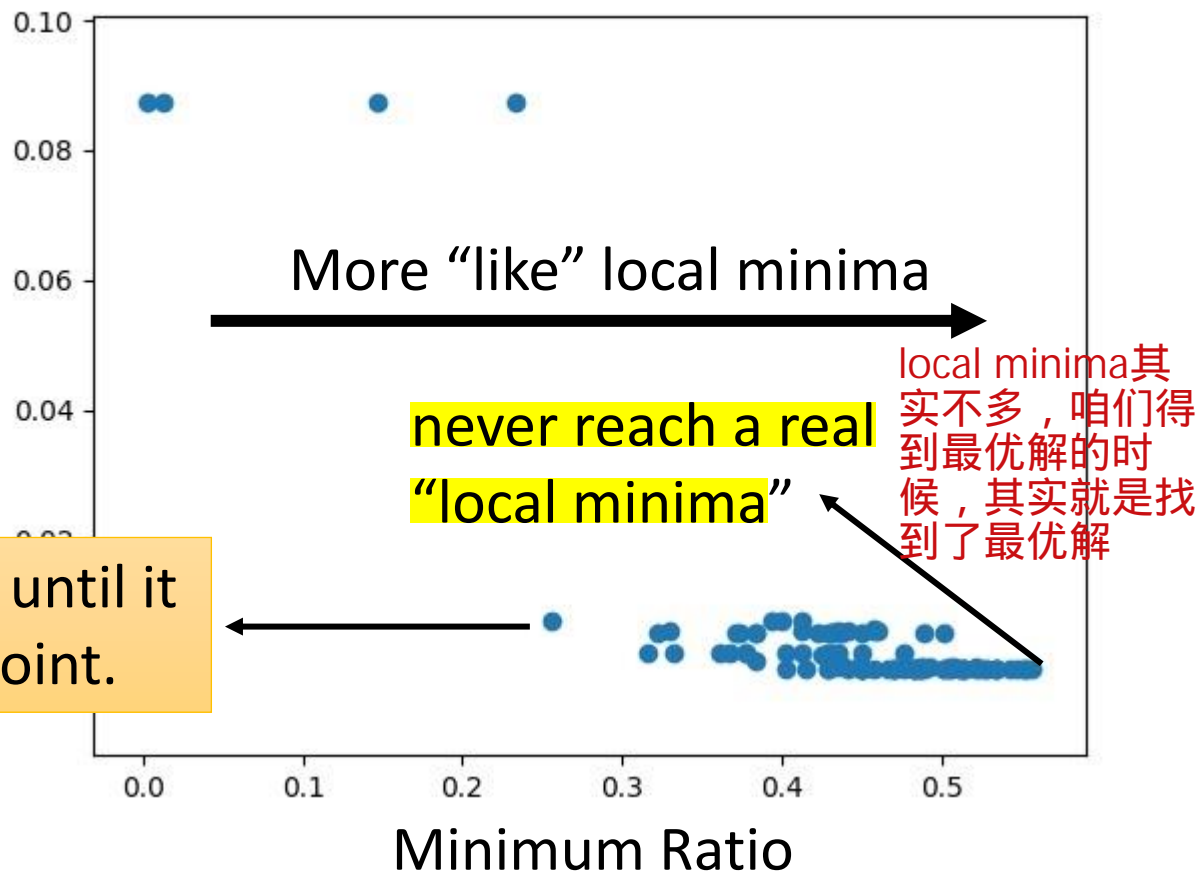
Saddle point in
higher dimension?

When you have lots of parameters, perhaps local minima is rare?

Empirical Study


Training
Loss

Train a network once, until it
converges to critical point.



$$\text{Minimum ratio} = \frac{\text{Number of **Positive** Eigen values}}{\text{Number of Eigen values}}$$

Small Gradient ...



Tips for training: Batch and Momentum



Batch

批次，也叫mini-batch

Review: Optimization with Batch

$$\theta^* = \arg \min_{\theta} L$$

➤ (Randomly) Pick initial values θ^0

➤ Compute gradient $g^0 = \nabla L^1(\theta^0)$


update $\theta^1 \leftarrow \theta^0 - \eta g^0$

➤ Compute gradient $g^1 = \nabla L^2(\theta^1)$

update $\theta^2 \leftarrow \theta^1 - \eta g^1$

➤ Compute gradient $g^3 = \nabla L^3(\theta^2)$

update $\theta^3 \leftarrow \theta^2 - \eta g^3$



1 **epoch** = see all the batches once → **Shuffle** after each epoch

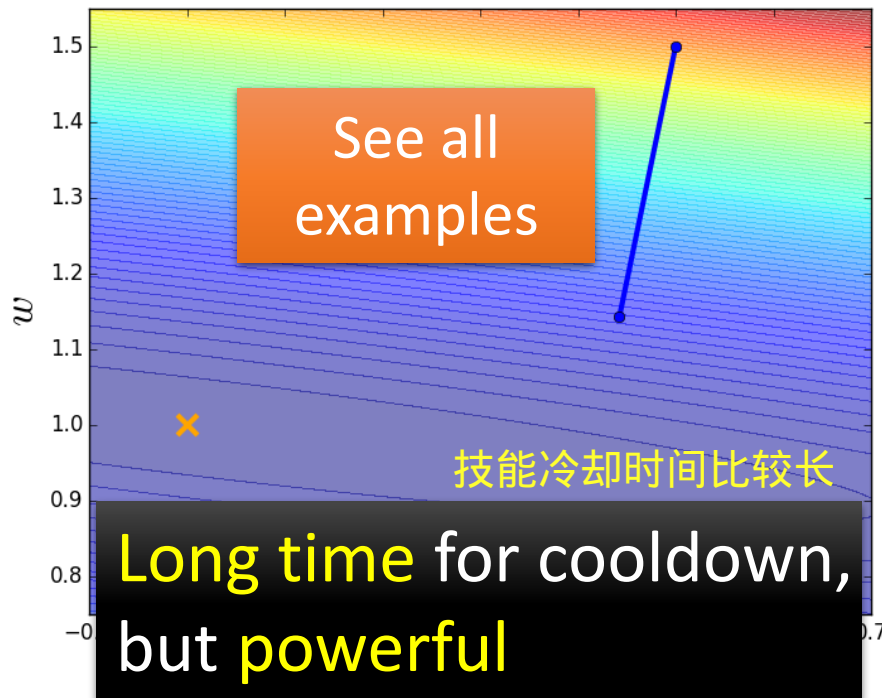
Small Batch v.s. Large Batch

Consider 20 examples ($N=20$)

Batch size = N (Full batch)

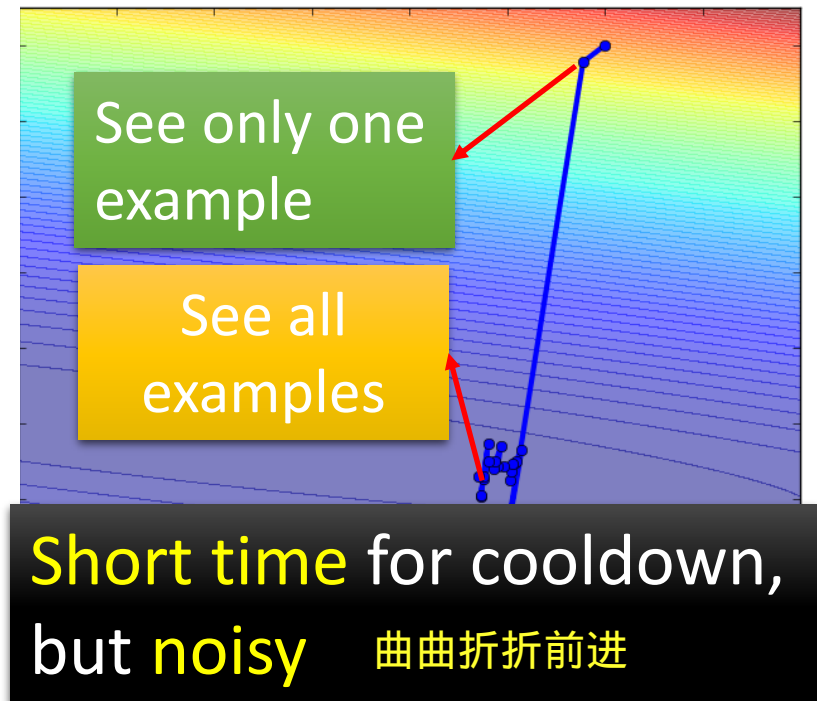
Full batch就是没用mini-batch

Update after seeing all
the 20 examples



Batch size = 1

Update for each example
Update 20 times in an epoch



Small Batch v.s. Large Batch


- Larger batch size does **not** require longer time to compute gradient (unless batch size is too large)

**Time for
each update**

MNIST: digit
classification

Tesla V100 GPU


可以做平行运算



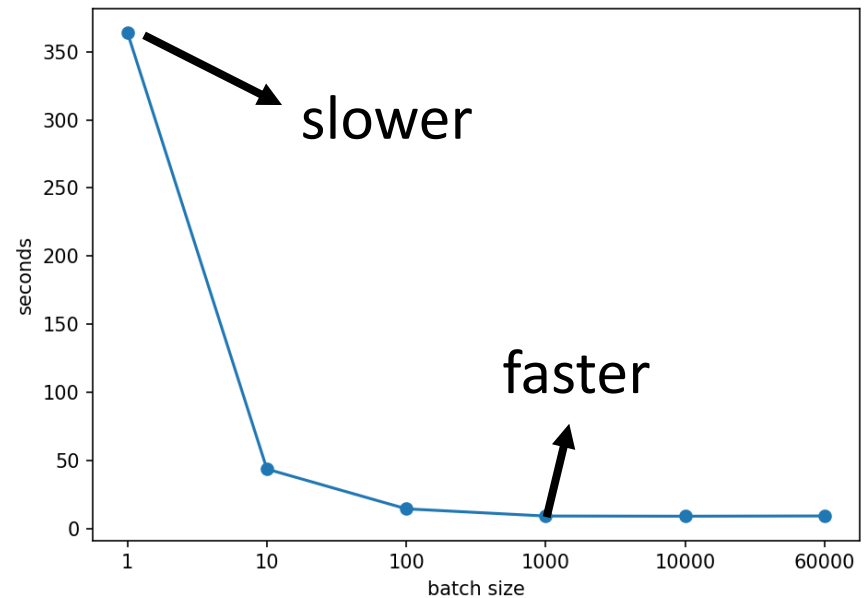
Small Batch v.s. Large Batch

- Smaller batch requires longer time for one epoch (longer time for seeing all data once)

Time for one **update**



Time for one **epoch**



Batch size越大，每epoch
需要的时间更短

60000 updates in one epoch


根据前面的实验，cd时长没差啦。

Small Batch v.s. Large Batch

Consider 20 examples ($N=20$)


Batch size = N (Full Batch)

Update after seeing all
the 20 examples



Batch size = 1

Update for each example
Update 20 times in an epoch




noisy有利于gradient descent

Small Batch v.s. Large Batch


MNIST

Accuracy vs. Batch Size



CIFAR-10



Accuracy vs. Batch Size



- Smaller batch size has better performance
- What's wrong with large batch size? **Optimization** Fails

Small Batch v.s. Large Batch

- Smaller batch size has better performance
- “Noisy” update is better for training



第二次epoch的时候，
不会stuck在第一epoch的地方，从而
继续更新参数，优化模型

Small Batch v.s. Large Batch

- Small batch is better on testing data?

SB = 256

LB =


0.1 x data set

Name	Network Type	Data set
F_1	Fully Connected	MNIST (LeCun et al., 1998a)
F_2	Fully Connected	TIMIT (Garofolo et al., 1993)
C_1	(Shallow) Convolutional	CIFAR-10 (Krizhevsky & Hinton, 2009)
C_2	(Deep) Convolutional	CIFAR-10
C_3	(Shallow) Convolutional	CIFAR-100 (Krizhevsky & Hinton, 2009)
C_4	(Deep) Convolutional	CIFAR-100




Name	Training Accuracy		Testing Accuracy	
	SB	LB	SB	LB
F_1	99.66% \pm 0.05%	99.92% \pm 0.01%	98.03% \pm 0.07%	97.81% \pm 0.07%
F_2	99.99% \pm 0.03%	98.35% \pm 2.08%	64.02% \pm 0.2%	59.45% \pm 1.05%
C_1	99.89% \pm 0.02%	99.66% \pm 0.2%	80.04% \pm 0.12%	77.26% \pm 0.42%
C_2	99.99% \pm 0.04%	99.99% \pm 0.01%	89.24% \pm 0.12%	87.26% \pm 0.07%
C_3	99.56% \pm 0.44%	99.88% \pm 0.30%	49.58% \pm 0.39%	46.45% \pm 0.43%
C_4	99.10% \pm 1.23%	99.57% \pm 1.84%	63.08% \pm 0.5%	57.81% \pm 0.17%

Small Batch v.s. Large Batch

- Small batch is better on testing data?



Small Batch v.s. Large Batch

	Small	Large
Speed for one update (no parallel)	Faster	Slower
Speed for one update (with parallel)	Same	Same (not too large)
Time for one epoch	Slower	Faster 
Gradient	Noisy	Stable
Optimization	Better 	Worse
Generalization	Better 	Worse

Batch size is a hyperparameter you have to decide.

Have both fish and bear's paws?

- Large Batch Optimization for Deep Learning: Training BERT in 76 minutes (<https://arxiv.org/abs/1904.00962>)
- Extremely Large Minibatch SGD: Training ResNet-50 on ImageNet in 15 Minutes (<https://arxiv.org/abs/1711.04325>)
- Stochastic Weight Averaging in Parallel: Large-Batch Training That Generalizes Well (<https://arxiv.org/abs/2001.02312>)
- Large Batch Training of Convolutional Networks (<https://arxiv.org/abs/1708.03888>)
- Accurate, large minibatch sgd: Training imagenet in 1 hour (<https://arxiv.org/abs/1706.02677>)

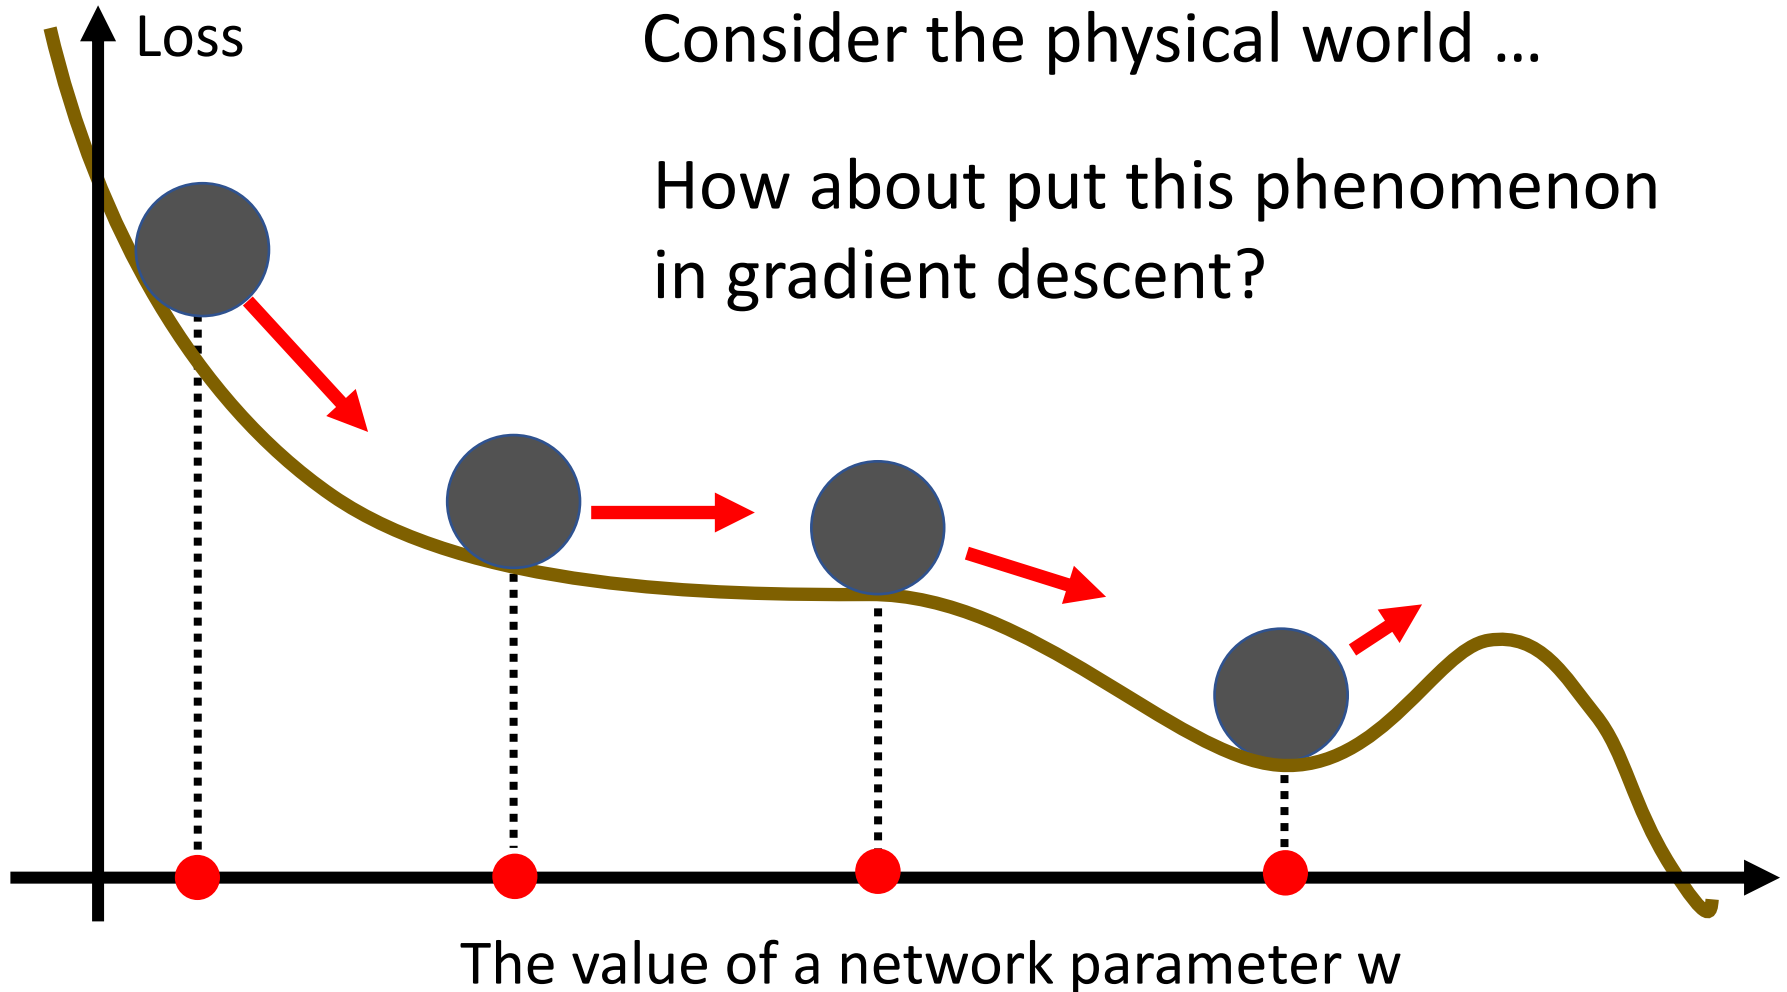
Momentum

动量

Small Gradient ...


Consider the physical world ...

How about put this phenomenon
in gradient descent?



(Vanilla) Gradient Descent

adj. 一般的



Starting at θ^0

Compute gradient g^0

Move to $\theta^1 = \theta^0 - \eta g^0$


Compute gradient g^1

Move to $\theta^2 = \theta^1 - \eta g^1$

⋮

Gradient Descent + Momentum

Movement: **movement of last step** minus **gradient** at present



Starting at θ^0 第0步

Movement $m^0 = 0$

Compute gradient g^0

Movement $m^1 = \lambda m^0 - \eta g^0$

Move to $\theta^1 = \theta^0 + m^1$

Compute gradient g^1 第1步

Movement $m^2 = \lambda m^1 - \eta g^1$

Move to $\theta^2 = \theta^1 + m^2$ 第2步

Movement not just based on gradient, but previous movement.

Gradient Descent + Momentum

Movement: **movement of last step** minus **gradient** at present

\mathbf{m}^i is the weighted sum of all the previous gradient: $\mathbf{g}^0, \mathbf{g}^1, \dots, \mathbf{g}^{i-1}$

$$\mathbf{m}^0 = \mathbf{0}$$

$$\mathbf{m}^1 = -\eta \mathbf{g}^0$$

$$\mathbf{m}^2 = -\lambda \eta \mathbf{g}^0 - \eta \mathbf{g}^1$$

\vdots

Starting at $\boldsymbol{\theta}^0$

Movement $\mathbf{m}^0 = \mathbf{0}$

Compute gradient \mathbf{g}^0

Movement $\mathbf{m}^1 = \lambda \mathbf{m}^0 - \eta \mathbf{g}^0$

Move to $\boldsymbol{\theta}^1 = \boldsymbol{\theta}^0 + \mathbf{m}^1$


Compute gradient \mathbf{g}^1

Movement $\mathbf{m}^2 = \lambda \mathbf{m}^1 - \eta \mathbf{g}^1$

Move to $\boldsymbol{\theta}^2 = \boldsymbol{\theta}^1 + \mathbf{m}^2$

Movement not just based on gradient, but previous movement.

Gradient Descent + Momentum



Concluding Remarks

- Critical points have zero gradients. $\nabla f = 0$ critical point
- Critical points can be either saddle points or local minima.
 - Can be determined by the Hessian matrix.
 - It is possible to escape saddle points along the direction of eigenvectors of the Hessian matrix.
 - Local minima may be rare.
- **Smaller batch size** and momentum help escape critical points.

Acknowledgement

- 感謝作業二助教團隊(陳宣叡、施貽仁、孟妍李威緒)幫忙跑實驗以及蒐集資料